



Investor Pack

The WhatsApp employee for Lagos & Abuja businesses.

Section	Source
1. Investor brief	INVESTOR.md
2. Pricing & packaging	PRICING.md
3. Unit economics	ECONOMICS.md

Compiled 26 May 2026 · Yori Ajagun · yoriajagun08@gmail.com

§ Investor brief

ReachNG — Investor Brief

The agentic employee for premium Nigerian SMEs. Built in Lagos. Operating in Lagos and Abuja.

> ReachNG is the WhatsApp-native customer operator for premium Lagos and Abuja SMEs. EYO drafts every reply in the owner's voice, parses GTBank/OPay transfer screenshots, qualifies leads, reactivates dormant customers, and briefs the founder every morning. Every message is human-approved by default. We don't replace the team — we give them a colleague that never sleeps.

Last updated: 2026-05-20 · Stage: Pre-revenue, live product, first paid client imminent

1 · One-page snapshot

What we sell	Monthly SaaS subscription. Owner pairs ReachNG to their existing WhatsApp number via Unipile QR. EYO operates from that number under HITL approval.
Who we sell to	Premium SMEs in Lagos and Abuja: luxury hospitality, real estate agencies (Banana Island / Ikoyi / Maitama tier), clinics, commercial law, family offices, professional services.
Pricing (live)	Starter ■150k / Growth ■300k / Scale ■600k per month. 15% off annual prepay.
Per-client gross margin	66 – 94% across all profile × tier combinations (see §3)
Stack	Python 3.12 + FastAPI + Mongo Atlas + Anthropic Haiku 4.5 + OpenAI Whisper + Unipile + Resend + Paystack + PostHog. Railway deploy.
HQ	Lagos, Nigeria
Founder	Oluyori Ajagun — back-office and risk operator (WFunded, formerly). Built the product from inside the pain it solves.
Capital required for first 10 paying clients	Negligible. Free tiers cover all infra. Cash positive in month 1 at first paid client.

2 · The market

TAM — Premium Lagos + Abuja SMEs

- **Lagos:** ~2.5M registered small/medium businesses. Premium subset (>20 staff, >■50M revenue, or luxury-coded sector) ≈ **50,000-80,000 businesses**.
- **Abuja:** ~500,000 businesses. Premium subset (Maitama, Wuse, Asokoro, Jabi) ≈ **10,000-15,000**.

- **Combined premium tier = ~60,000-95,000 businesses.**

Beachhead — Our ICP within TAM

Premium businesses that already live on WhatsApp for inbound, have 5+ staff, transact via bank transfer, and feel the late-evening leakage problem acutely:

- Luxury hospitality (premium restaurants, hotels, event venues, lounges)
- Real estate (sales + rental, Banana Island / Ikoyi / Eko Atlantic / Maitama tier)
- Clinics (aesthetic, dental, wellness, fertility)
- Commercial law and advisory
- Family offices
- High-ticket retail (auto, jewellery, art)

Conservative ICP count: **~5,000-10,000 businesses across both cities**. At 2% penetration = 100-200 clients = **■18-36M MRR** on the current ladder.

Why Lagos + Abuja, not all-Nigeria

- Concentration of premium businesses
- WhatsApp Business penetration ≈ 95%
- Naira + Paystack-native, no FX friction
- Founders are decision-makers who can sign in days, not months
- Operator (Yori) is on the ground in Lagos with a network

National expansion (Port Harcourt, Ibadan, Kano) is a year-2 move, not year-1.

Buy-trigger

The premium-SME version of "AI for business" tools is broken in Nigeria. Generic foreign SaaS (HubSpot, Drift, Intercom) doesn't speak Paystack, naira, or Pidgin. Local tools are either thin CRMs or bulk SMS senders. There's no agentic WhatsApp operator built for this market. We're the first.

3 · Unit economics (see [ECONOMICS.md](#) for the full breakdown)

Cost floor per client per month (Unipile-paired, Mongo on free tier today)

Profile	Per-call AI	Per-client fixed	Platform share (÷10)	Total cost
Light	■2,520	■25,600	■6,100	■34,200
Medium	■7,620	■25,600	■6,100	■39,300
Heavy	■18,840	■25,600	■6,100	■50,500

Light = small restaurant / single agent. Medium = premium hospitality / mid clinic / established agency. Heavy = high-volume clinic / multi-location restaurant / busy agency.

Gross margin per tier × profile

	Light client	Medium client	Heavy client
Starter ■150k	77%	74%	66%
Growth ■300k	89%	87%	83%
Scale ■600k	94%	93%	92%

Worst case is 66% gross margin (Heavy-usage client on Starter plan). Best case is 94%.

What scales how

- **Per-call costs** (Whisper, Haiku, Vision): scale linearly with usage. Buffered ~30-50% above raw vendor pricing in the ledger to absorb FX swings and retries.
- **Per-client fixed:** Unipile per-account fee (~■24k) is the dominant cost. Scales linearly with client count.
- **Platform fixed:** Railway + Resend + PostHog + Mongo are roughly fixed regardless of client count, but Mongo upgrades from free to M10 (~■91k/mo) at ~10-12 clients. Modelled below.

4 · Revenue projections — three scenarios

All projections use Ladder B (■150 / 300 / 600). Pilot pricing (33% off first 3 clients × 90 days) factored into year-1 numbers.

Mix assumption

- Starter clients: ~50% of book (smaller premium SMEs)
- Growth clients: ~40% of book (mid-size premium)
- Scale clients: ~10% of book (luxury / multi-location)

This is the conservative weighting — Lagos/Abuja market data suggests Growth tier may end up ~50% as buyers self-select up.

Scenario A — Conservative (slow ramp)

Month	Clients	MRR	Monthly cost	Monthly profit	Margin
M3	3	■450,000	■170,000	■280,000	62% (pilot pricing)
M6	8	■1,800,000	■330,000	■1,470,000	82%
M12	20	■5,250,000	■830,000*	■4,420,000	84%
M24	50	■13,125,000	■2,080,000	■11,045,000	84%

*Includes Mongo M10 transition around client 10.

Year-1 ARR exit: ■63M. Year-2 ARR exit: ■157M.

Scenario B — Base case

Month	Clients	MRR	Monthly cost	Monthly profit	Margin
M3	5	■750,000	■210,000	■540,000	72% (pilot pricing)
M6	15	■3,375,000	■620,000	■2,755,000	82%
M12	40	■10,500,000	■1,720,000	■8,780,000	84%
M24	100	■26,250,000	■4,170,000	■22,080,000	84%

Year-1 ARR exit: ■126M. Year-2 ARR exit: ■315M.

Scenario C — Aggressive (network effects + referral kick in)

Month	Clients	MRR	Monthly cost	Monthly profit	Margin
M3	8	■1,200,000	■330,000	■870,000	73% (pilot pricing)
M6	25	■5,625,000	■1,040,000	■4,585,000	82%
M12	75	■19,687,500	■3,200,000	■16,487,500	84%
M24	200	■52,500,000	■8,340,000	■44,160,000	84%

Year-1 ARR exit: ■236M. Year-2 ARR exit: ■630M.

Margin trajectory (all scenarios)

Margins **expand with scale**, not contract:

- Sub-10 clients: 66-89% (pilot pricing + lower volume)
- 10-50 clients: 82-87% (Mongo M10 cost absorbed)
- 50+ clients: 84-89% (per-client fixed share drops further)

The Unipile per-account fee is the only line that scales linearly with clients. Everything else compounds in our favour.

4.5 · Road to 1,000 clients (the headline goal)

1,000 paying clients ≈ 10-20% of the addressable premium SME beachhead in Lagos and Abuja. The path, broken into quarterly milestones, all on Ladder B pricing:

Year 1 — Foundation (0 → 100)

Quarter	Target clients (cumulative)	Net adds	MRR at quarter-end	Notes
Q1 (M1-3)	5	+5	■750,000	Pilot batch, hand-onboarded, 33% pilot discount in effect

Quarter	Target clients (cumulative)	Net adds	MRR at quarter-end	Notes
Q2 (M4-6)	25	+20	■5,625,000	Pricing live, founder-led sales, referrals start
Q3 (M7-9)	60	+35	■13,500,000	Vertical case studies live, Twitter authority cadence
Q4 (M10-12)	100	+40	■22,500,000	First sales/CS hire mid-Q4

Year-1 exit: 100 clients · ■22.5M MRR · ■270M ARR

Year 2 — Velocity (100 → 600)

Quarter	Target clients	Net adds	MRR	Notes
Q5 (M13-15)	175	+75	■39.4M	First SDR hire, paid acquisition unlocked
Q6 (M16-18)	275	+100	■61.9M	Agency partnership channel live
Q7 (M19-21)	425	+150	■95.6M	Conference + brand campaign, Abuja office
Q8 (M22-24)	600	+175	■135.0M	2nd engineer hire

Year-2 exit: 600 clients · ■135M MRR · ■1.62B ARR

Year 3 — Compounding (600 → 1,000)

Quarter	Target clients	Net adds	MRR	Notes
Q9 (M25-27)	750	+150	■168.8M	Voice Operator (outbound) launches
Q10 (M28-30)	850	+100	■191.3M	National expansion (Port Harcourt, Ibadan)
Q11 (M31-33)	925	+75	■208.1M	Vertical SDK / API offering
Q12 (M34-36)	1,000	+75	■225.0M	West Africa pilot (Accra, Nairobi)

Year-3 exit at 1,000 clients: ■225M MRR · ■2.7B ARR · ~85% blended margin · ~■191M monthly net profit.

Cost at each milestone

Includes Claude Pro/Max operator-tooling cost in the founder-toolset line (Pro at start, Max 5× at ~15 clients, Max 20× at ~50 clients).

Clients	MRR	Direct cost	Team cost	Mongo tier	Operator tools	Total cost	Net margin
25	■5.6M	■820k	Founder only	Atlas free	Claude Max 5× ■160k	■980k	83%
100	■22.5M	■3.5M	+1 ops/CS	M10 ■91k	Max 20× ■320k	■4.4M	80%
250	■56.3M	■8.5M	+SDR + a ccounant	M20 ■232k	Max 20× ■320k	■11.9M	79%
500	■112.5M	■16.7M	+2 engineers + CS	M20	2× Max 20× ■640k	■24.7M	78%
1,000	■225M	■32M	~15-person org	M30 ■617k	4× Max 20× ■1.28M	■49.3M	78%

Read: even at 1,000 clients with a 15-person Lagos team and full Claude Max for every operator, blended net margin sits at ~78%. The model is built for compounding profitability, not just compounding revenue.

Hire plan along the curve

Milestone	Hires added
25 clients	First ops / customer success generalist
75 clients	First SDR (replaces founder-led outreach)
100 clients	Accountant / bookkeeper (PT)
250 clients	Lead engineer + CS lead
500 clients	2 more engineers + 2 more SDR + Head of Sales
1,000 clients	Country leads (Lagos / Abuja / PH), VP Eng, VP Customer Ops

Final-state org: ~15 people. Revenue per employee at 1,000 clients = ■180M/year. Industry benchmark for high-margin SaaS is ■40-80M/employee. We'd be top-quartile.

What the 1,000-client outcome looks like

- ■2.7B annualised revenue

- **■2.3B annualised net profit**
 - **15-person team across Lagos / Abuja / Port Harcourt**
 - **Default agentic-employee platform for the Nigerian premium SME tier**
 - **Optionality:** vertical SDKs (real estate, hospitality), API offering, regional expansion, acquisition target for global CRM consolidators (HubSpot, Salesforce) at 5-8× ARR multiple = **■13-22B valuation**
-

5 · Customer acquisition

Channel today

- **Own SDR funnel** (Google Maps Places + Apollo + Apify + Twitter/LinkedIn signals → enrichment → Haiku-drafted email/WhatsApp outreach → HITL approval by Yori → send from hello@reachng.ng).
- **Pilot Application waitlist** at reachng.ng/waitlist (with PostHog tracking for source attribution).
- **Founder authority** — Yori posting on Twitter/LinkedIn about agent operations as the proof.

CAC

- Variable cost: ~**■8k/mo** Apify token, ~**■5k/mo** for Resend + Anthropic on outreach. Call it **■15k/mo of variable acquisition cost**.
- Yori's time: a real cost but uncapitalised for first ~20 clients.
- **Pure CAC per client** ≤ **■5,000** through bootstrapped channels. Effectively negligible.

LTV (conservative)

- Average contract: 24 months (premium SaaS norm; switching cost is high once WhatsApp paired)
- Blended ARPC: **■250,000/mo** (across Starter/Growth/Scale)
- Gross margin: 84% blended
- **LTV = 24 × 250,000 × 0.84 = ■5,040,000 per client**

LTV / CAC

At **■5,000 CAC** and **■5.04M LTV** = **1,008× LTV/CAC ratio**. The benchmark anyone respects is >3×. This number won't hold once we run paid acquisition or hire SDRs, but the bootstrapped phase is exceptional.

Paid CAC ceiling

If LTV is **■5M** and we want to keep LTV/CAC > 5×, we can spend up to **■1M acquiring a client** before economics break. Plenty of headroom for paid social, partnership commissions, conference sponsorships when the time comes.

6 · Why this is defensible (the moats — see also memory/project_reachng_ai_moats.md)

The Claude / OpenAI / Google horizontal AI commoditisation question matters. Five moats:

- **Distribution, not intelligence.** The hard part isn't drafting a reply — Haiku does that. The hard part is being wired into the SME's *actual* WhatsApp number, with HITL queue, receipt parsing, ledger

reconciliation, Paystack subscription billing. Generic Claude doesn't ship that.

- **Integration tax.** Unipile + Meta + Anthropic + OpenAI + Resend + Paystack + Mongo + LeanScrape took 4 months to integrate cleanly. Each week of work is a week a competitor would need.
- **Vertical memory.** PAYE bands, Lagos Tenancy Law escalation, GTB/OPay/Kuda transfer patterns, Pidgin grammar, "Thank you sir" close — all encoded. A horizontal AI has to be told every time.
- **The HITL UX is the product.** Owners want draft-and-tap, not autonomous AI. That's a UX product, not a model capability. Anthropic isn't shipping it.
- **Closed feedback loop per client.** Every approved/edited draft tunes the next one in that client's account. T0.4 outcome learning compounds per-client.

When Haiku gets 10× better, ReachNG gets 10× better — same code, cheaper inference. The competitive layer is everything *around* the model.

7 · The voice-note safe-switch (a representative moat)

A specific example of how we win on trust where generic AI loses.

Lagos WhatsApp customers send voice notes in five languages: English, Pidgin, Yoruba, Igbo, Hausa. OpenAI Whisper handles English and Pidgin well. Yoruba/Igbo/Hausa work for clean sentences but slip on idioms or colloquial speech.

A generic AI tool would confidently mistranslate. ReachNG enforces a strict contract:

- **English / Pidgin:** transcript flows to drafter as text.
- **Yoruba/Igbo/Hausa, high confidence (≥90%):** transcript + Haiku-powered English translation surfaced side-by-side. EYO drafts the reply in English.
- **Anything low confidence: EYO does not draft.** HITL banner reads "language uncertain, listen to the audio before sending anything."

The line we won't cross: never let EYO confidently guess at a voice note it didn't fully understand. Premium owners pay extra for that honesty because their brand is on the line.

8 · Risk register

Top five risks, ranked by impact × probability.

1 · Unipile session expiry (mitigation in flight)

WhatsApp linked-device sessions expire after ~14 days of phone inactivity. Silent expiry = silent product stop. Mitigation: scheduled health-check (6h cadence) + portal banner + owner WhatsApp alert + PostHog event. Queued for build before first paid client. ~1 day work.

2 · Meta deprecates linked-device API or restricts third-party access

Lower-probability but high-impact. Mitigation: we already have Meta Cloud API integration as the fallback. If Unipile breaks, clients can migrate (with friction).

3 · Mongo Atlas free tier outgrown

Around 10-12 clients we move to M10 (■91k/mo). Modelled in projections. Not a surprise.

4 · Currency volatility

All revenue is in NGN, most vendor costs are USD (Anthropic, OpenAI, Unipile, Resend). 20% NGN devaluation = ~10% cost-side hit on per-call costs (we have 30-50% buffer in the ledger). Manageable through pricing reviews.

5 · NDPR / data regulation

Lawful basis already encoded into BYO Leads ingest (consent gates, source-of-consent attestation, audit log, hard-delete on request). Lagos lawyer review queued before first paid client signs MSA.

9 · Use of capital (if any is raised)

Not raising on this doc, but plausible uses:

- **■5-20M pre-seed friends-and-family:** 6 months of personal runway for Yori to focus full-time, ~■2M Apify/ads budget, modest hardware/AV for video demos.
- **■50-150M seed (US\$ equivalent at growth):** First sales/SDR hire (Lagos), brand campaign for premium positioning, lawyer + accounting retainer, Mongo Atlas dedicated cluster, customer success ops, expand to Port Harcourt + Ibadan.
- **■500M+ A round:** national + West African expansion, agency partnership channel, RHEL infrastructure scaling, voice-operator outbound build (Phase 5 in PLAN.md).

The honest position: the unit economics are strong enough to bootstrap to ■150M+ ARR without external capital. Raising would accelerate, not enable.

10 · Why now

- **WhatsApp Business penetration in Nigeria** crossed 95% in 2024-2025. The channel is universal.
 - **Haiku 4.5 / GPT-4o latency** crossed the "feels live" threshold in 2025 — sub-2s drafts are now possible at ~■4 per call.
 - **Lagos premium SMEs** are exiting COVID-era cost discipline and looking for force-multipliers, not headcount.
 - **Founder-led AI tools** are out-shipping enterprise SaaS for the SME tier worldwide; Nigeria has no homegrown answer yet.
 - **The integration tax** that protects us only gets thicker with every API we add. The moat compounds while competitors evaluate the entry cost.
-

11 · Status check

- **■ Product live** at reachng.ng
- **■ Self-serve signup + Paystack billing** wired (end-to-end tested)
- **■ HITL queue + WhatsApp pairing flow** shipped

- ■ Voice + Receipt + Morning Brief features shipped (T0.1, T0.2)
 - ■ PostHog product analytics live
 - ■ Pricing locked at Ladder B (■150 / 300 / 600)
 - ■ First paid client target: imminent
 - ■ Unipile session-expiry health loop: queued
 - ■ Phase 1.6 client book onboarding (vCard / WhatsApp share / paste): queued
 - ■ Lawyer-reviewed MSA + DPA + NDA: drafted, pending Lagos counsel review
-

12 · How to verify

This is not a slide deck. Every number above can be verified against:

- ECONOMICS.md — per-call cost rates with vendor cross-check
 - PRICING.md — three pricing ladders modelled
 - `services/usage_meter.py` — live cost ledger code
 - `api/billing.py` + `templates/dashboard.html` — operator visibility today
 - Live admin dashboard at `/dashboard` (Basic Auth gated)
 - Public landing + signup at reachng.ng
-

Contact: Oluyori Ajagun · hello@reachng.ng · WhatsApp +234 816 458 3657

§ Pricing & packaging

ReachNG — Pricing Working Doc

Locked premises (don't relitigate):

- **Unipile 100%** for WhatsApp. No Meta Cloud API. Per-client fixed cost ~\$24k/mo + shared infra.
- **HITL non-negotiable.** Every outbound is human-approved unless owner explicitly flips Autopilot per reply-type.
- **Lagos + Abuja pilot.** Premium SME positioning. Hand-onboarded, selective intake.
- **First paid client target:** lock pricing today so signup + pricing page reflect reality before any paid signup.

Last updated: 2026-05-20

1 · Cost floor per client (Unipile-only, Mongo Atlas free tier)

From [ECONOMICS.md](#) \$6, with Meta Cloud option removed, Unipile assumed, and Mongo on free tier (\$0 today):

Profile	Per-call AI	Per-client fixed (Unipile + Resend)	Platform share (÷10 clients)	Total cost/mo
Light	2,520	12,000	6,100	~\$20,620
Medium	7,620	12,000	6,100	~\$25,720
Heavy	18,840	12,000	6,100	~\$36,940

These are the floors. Any price below means a loss on that profile.

Unipile pricing reality (corrected 2026-05-20): First 10 connected accounts = \$55/mo flat (~\$88,000 total = \$8,800/client). Then \$5.5 per additional (11-50), \$5 (51-200), \$4.5 (201-1k). Per-client Unipile cost stays in the \$7,400-8,800 band across all scales.

Platform share (\$6,100/client at 10 clients) includes Railway + Resend Pro + PostHog + Claude Pro subscription + domain + Apify SDR funnel. Claude upgrades to Max 5x (~\$160k) around 15 clients, Max 20x (~\$320k) around 50 clients. Still negligible per-client at scale (~\$3k each at 100 clients).

Note on Mongo: you're on Atlas free tier (M0, 512MB) today, so MongoDB cost is \$0. When client volume + memory data outgrows free tier (around client 8-12 depending on usage), expect to move to M10 (~\$91,000/mo). Re-run the floor table at that point — every profile will rise by ~\$9k. Still healthy on the locked ladder.

2 · The three plausible ladders

Ladder A — Light bump (\$100 / 200 / 400)

Plan	Price	Margin on Light	Margin on Medium	Margin on Heavy
Starter	\$100,000	+ \$62,600 (63%)	+ \$57,500 (58%)	+ \$46,300 (46%)

Plan	Price	Margin on Light	Margin on Medium	Margin on Heavy
Growth	■200,000	+■162,600 (81%)	+■157,500 (79%)	+■146,300 (73%)
Scale	■400,000	+■362,600 (91%)	+■357,500 (89%)	+■346,300 (87%)

Read: Starter tight on heavy clients (46%). Growth and Scale clean.

Ladder B — Operator-grade (■150 / 300 / 600) — *previously locked, superseded* — see note below

Plan	Price	Margin on Light	Margin on Medium	Margin on Heavy
Starter	■150,000	+■112,600 (75%)	+■107,500 (72%)	+■96,300 (64%)
Growth	■300,000	+■262,600 (88%)	+■257,500 (86%)	+■246,300 (82%)
Scale	■600,000	+■562,600 (94%)	+■557,500 (93%)	+■546,300 (91%)

Read: every tier × every profile margin >64%. Worst case (Heavy on Starter) is still healthy 64%. Headroom for surprise spikes and the eventual Mongo M10 + Claude Max upgrades.

Ladder C — Premium-anchor (■200 / 400 / 800)

Plan	Price	Margin on Light	Margin on Medium	Margin on Heavy
Starter	■200,000	+■162,600 (81%)	+■157,500 (79%)	+■146,300 (73%)
Growth	■400,000	+■362,600 (91%)	+■357,500 (89%)	+■346,300 (87%)
Scale	■800,000	+■762,600 (95%)	+■757,500 (95%)	+■746,300 (93%)

Read: aggressive but defensible. "■200k/mo is two days of a senior receptionist's salary, except this one never sleeps." Risk: prices out non-luxury verticals (beauty, small hospitality), narrows TAM.

3 · How the math behaves at scale

At 10 paying clients (Ladder B, mix of usage):

Mix	Monthly revenue	Monthly cost	Monthly profit	Margin
5 Starter + 3 Growth + 2 Scale	■2.85M	■420,000	■2.43M	85%
3 Starter + 5 Growth + 2 Scale	■3.15M	■425,000	■2.73M	87%
2 Starter + 4 Growth + 4 Scale	■3.90M	■440,000	■3.46M	89%

Read: annualised ■29-42M revenue from 10 clients. Strong unit economics. Even on Ladder B alone, 10 clients gives a healthy runway with very thin OpEx (Atlas free tier + Railway + Resend = real fixed costs under ■60k/mo).

Reaching 50 clients on Ladder B

Assume 25 Starter, 20 Growth, 5 Scale = ■12.75M MRR = ■153M ARR.

■■ At ~10-12 paying clients we'll outgrow Atlas free tier. Move to M10 (~■91k/mo). Blended margin still ~88% at scale. Real numbers.

3.5 · Bundle plan (Premium only)

For owners running multiple businesses or multi-location operations. **Bundle is exclusive to Premium** — Starter and Growth do not offer add-on accounts. This keeps tier ladders clean and forces the right buyer into the right plan: if you have 2+ WhatsApp lines worth running on EYO, you're on Premium.

Pricing

- **Premium base:** ■250,000/mo (founder) · ■300,000/mo (regular) — includes 1 paired WhatsApp number + 1 vertical playbook
- **Each additional WhatsApp number:** +■125,000/mo (founder) · ■150,000/mo (regular). Half the base price, per the simple rule.
- Each additional number can be in **any vertical** — EYO loads the matching per-vertical playbook, tone profile, and Control Room dashboard for that number. A single owner can run a restaurant + RE agency + clinic under one bundle.

Example math (founder pricing)

Lagos owner with 3 businesses across verticals:

Operation	Vertical	Founder price
Altitude Lagos (rooftop restaurant)	hospitality	■250,000
+ Ikoyi Prime (real estate agency)	real_estate	■125,000
+ Lekki Aesthetic Clinic	clinics	■125,000
Bundle total		■500,000/mo
(3 standalone Premium subscriptions would be)		■750,000/mo
Owner saves		■250,000/mo

Margin on each add-on account

Incremental cost per additional account ≈ ■15,000 (Unipile add-on ~■8,800 + small Resend/Mongo share + per-call delta).

- Add-on price ■125,000 – ■15,000 cost = **■110,000 profit per additional (88% margin)**

Stronger margin than the base Premium tier itself, because the platform overhead is already paid for by the first seat.

Billing implementation (Phase 2 — manual today)

Today: founder-led close. Owner signs up via standard Premium signup → Paystack charges █250k for first month → after onboarding call, we send a separate Paystack invoice for each additional number → manual reconciliation.

Phase 2 (once 3+ bundle clients close): add a "+ Additional Line" UI in the client portal that auto-charges █125k via Paystack Subscriptions and provisions a second WhatsApp pairing slot.

4 · Annual prepay (keep or adjust)

Current PRODUCTS.md offers **15% off annual prepay**. With Ladder B:

Plan	Monthly	Annual prepay (15% off)	Cash collected	Equivalent monthly
Starter	█150,000	█1,530,000	█1,530,000	█127,500
Growth	█300,000	█3,060,000	█3,060,000	█255,000
Scale	█600,000	█6,120,000	█6,120,000	█510,000

Cash on day 1 if a client takes Scale annual = █6.12M. That alone funds 4 months of platform fixed costs.

Recommendation: keep 15% annual discount. It's the founder-friendly close on a longer commitment.

5 · Plan inclusions (what changes by tier)

Before locking prices, lock what's IN each plan. Strawman:

Feature	Starter	Growth	Scale
EYO on owner's WhatsApp via Unipile pairing	✓	✓	✓
HITL draft queue	✓	✓	✓
Voice note transcription (EN/Pidgin/Yo/Ig/Ha safe-switch)	✓	✓	✓
Receipt screenshot reader	✓	✓	✓
Morning Owner Brief	✓	✓	✓
Business Brief intake + tone learning	✓	✓	✓
Per-vertical Control Room	✓	✓	✓

Feature	Starter	Growth	Scale
Re-engagement (Past / Dormant / Hot bucket triage)	—	✓	✓
BYO Leads ingest (vCard / paste / CSV / WhatsApp share)	—	✓	✓
Outcome learning loop (T0.4)	—	✓	✓
Proactive Intelligence (festival nudges, capacity, birthdays)	—	—	✓
Co-Pilot (owner asks EYO questions)	—	—	✓
Multi-location coordination	—	—	✓
Voice Operator (outbound voice agent)	—	—	✓
Priority response (15 min during business hours)	—	—	✓
Architecture tour + bespoke modules	—	—	✓

This forces an honest reason for each upgrade. Owners can see what they get by stepping up.

6 · Pilot-phase pricing (first 3-5 clients only)

Live pricing should be Ladder B. But for the **first 3 pilot clients**, a one-time pilot discount makes sense to reduce friction and generate case studies:

Plan	Live price	Pilot price (first 90 days)	Lock
Starter	■150,000	■100,000	converts to live price month 4
Growth	■300,000	■200,000	converts month 4
Scale	■600,000	■400,000	converts month 4

Frame to pilot owners: *"You're getting our first-customer rate. After 90 days you move to standard pricing or cancel, your choice."* That's defensible, not a discount-spiral.

7 · Recommendation to lock today

- **Ladder B (■150 / 300 / 600 monthly).** 65-89% margin across all client profiles. Headroom for surprises.
- **15% off annual prepay.** Keep. Closes long deals fast.
- **Plan inclusions per \$5.** Forces clean upgrade paths.
- **Pilot pricing per \$6** for first 3 clients only, time-boxed to 90 days.

If this looks right, downstream code updates needed:

- `services/platform_settings.py::get_plan_pricing()` → set new tier values
- `templates/marketing/pricing.html` → headline numbers, feature delta matrix
- `templates/marketing/signup.html` → plan picker labels
- `api/marketing.py::PLAN_PRICING` → fallback constants
- `PRODUCTS.md` → canonical tier definitions
- Welcome email + portal copy → reflect new amounts

I can ship all 6 in one bundle in ~30 minutes once you sign off on Ladder B (or pick a different ladder, or tune the numbers).

8 · The questions only you can answer

Before I ship anything:

- **Ladder A / B / C — which one?** Default is B.
- **Pilot pricing — yes or no?** Default yes for first 3 clients, off after that.
- **What's the highest you'd actually quote a Banana Island estate agency or a Maitama clinic with a straight face?** Calibrates whether Scale should be ■600k, ■800k, or higher.
- **What's the *lowest* you'd take from a small Surulere restaurant founder who really wants in?** Calibrates whether we keep Starter at ■150k or add a Lite ■100k.

Drop your four answers and I'll lock + ship.

§ Unit economics

ReachNG — Unit Economics

Honest, operator-grade breakdown of what one paying client actually costs us to run, and what the existing admin dashboard sees vs. what it misses.

Last updated: 2026-05-20 (pre-pricing session)

1 · Cost categories

We split costs into three buckets:

Bucket	Scales with	Example
Per-call (metered)	Each AI call a client triggers	Whisper transcribe, Haiku draft, Vision receipt
Per-client fixed	Each connected client, regardless of volume	Unipile per-number, Meta WhatsApp messaging fee, Resend share
Platform fixed	Total system, not per client	Railway hosting, Mongo Atlas, PostHog, domain

The meter (`services/usage_meter.py`) tracks bucket 1 only. Buckets 2 and 3 need manual attribution — see §5.

2 · Per-call costs (the meter sees these)

Rates currently encoded in `services/usage_meter.py::FEATURE_COSTS`:

Feature	■/call	Underlying call	What triggers it
voice	8	OpenAI Whisper (~30–90s audio)	Every inbound voice note
receipt	10	Anthropic Haiku 4.5 vision	Every inbound image (receipt screenshot)
classifier	2	Haiku, ~200-400 tok	Every inbound text classified for emotion / intent
drafter	4	Haiku, ~400-700 tok	Every reply EYO drafts (HITL queue)
memory	3	Haiku, structured extraction	Every inbound that bumps the client memory layer

Feature	■/call	Underlying call	What triggers it
copilot	8	Haiku planner + narrator	Each owner Co-Pilot question

These are anchors, not contracts. They were set conservatively before live data. Reconcile against actual API invoices monthly via `services/usage_meter.py` ledger.

Reality-check against vendor pricing (2026-05 Nigerian rates)

- **OpenAI Whisper:** \$0.006/minute audio = ■9.60/min @ ■1600/USD. Most voice notes ≤ 1 min → ■5-10. **Ledger rate of ■8 ≈ realistic.**
- **Haiku 4.5:** \$0.80 input / \$4 output per million tokens. A 600-token draft (350 in / 250 out) = \$0.0013 = ■2.10. **Ledger rate of ■4 ≈ 1.9× safety buffer.**
- **Haiku 4.5 vision (receipt):** ~1500 image tokens + 300 output = ~\$0.0024 = ■3.80. **Ledger rate of ■10 is generous** (justified — vision retries on poor receipts).
- **Memory extraction:** ~300 tok in / 100 tok out = ■1.20. Ledger ■3, buffered.
- **Co-pilot:** Two Haiku calls (planner + narrator). ~■5 actual. Ledger ■8 buffered.

Net: the ledger over-estimates by ~30–50% versus raw API cost. That buffer absorbs retries, error paths, and FX swings. **Don't tune down until 60+ days of real data confirms it.**

3 · Per-call cost by client usage profile

Three realistic Lagos/Abuja archetypes. Numbers are **monthly counts** of each feature event.

Profile A — Light (small restaurant, beauty salon, single agent)

Feature	Count/mo	■/call	Total ■/mo
voice	60	8	480
receipt	30	10	300
classifier	300	2	600
drafter	200	4	800
memory	100	3	300
copilot	5	8	40
Per-call total			■2,520/mo

Profile B — Medium (premium hospitality, mid-size clinic, established RE agency)

Feature	Count/mo	■/call	Total ■/mo
voice	200	8	1,600
receipt	80	10	800
classifier	900	2	1,800

Feature	Count/mo	■/call	Total ■/mo
drafter	600	4	2,400
memory	300	3	900
copilot	15	8	120
Per-call total			■7,620/mo

Profile C — Heavy (high-volume clinic, multi-location restaurant, busy agency)

Feature	Count/mo	■/call	Total ■/mo
voice	500	8	4,000
receipt	150	10	1,500
classifier	2,500	2	5,000
drafter	1,500	4	6,000
memory	700	3	2,100
copilot	30	8	240
Per-call total			■18,840/mo

4 · Per-client fixed costs (NOT in the meter today)

These attach to a client by the act of connecting them, even if they barely use the system. **This is where pricing surprises live.**

Line item	Cost basis	Per-client ■/mo
Unipile WhatsApp account	Real published tiered pricing: \$55/mo flat for first 10 accounts, then \$5.5/account (11-50), \$5 (51-200), \$4.5 (201-1k), \$4 (1k-5k), \$3.5 (5k+). Effective: ■7,400-8,800 per client across all scales.	~■8,800 at 10 clients
Meta Cloud API messaging fee	First 1,000 service conversations/mo free; then ~\$0.0094/conversation for Nigeria utility tier	0 to ~■3,000
Resend transactional email	Shared \$20/mo plan ÷ ~20 clients = ~\$1 attributable	~■1,600

Line item	Cost basis	Per-client ■/mo
MongoDB Atlas share	Currently on free tier M0 (512MB). ■0 today. Move to M10 (~\$57/mo) at ~10-12 clients = ~■9,100 then.	~■0 (today)
Anthropic + OpenAI floor	API key flat fees ≈ \$0 once usage-based; only fixed minimums during low usage	0
Estimated per-client fixed		~■12,000/mo (today, Mongo free, 10-client scale)

Critical: the Unipile per-account fee dominates. ■24k/client is the floor before EYO has done a single Haiku call.

Alternative: Meta Cloud API direct (skip Unipile)

We have Meta Cloud integration coded too. Cost shifts:

- Unipile fee → ■0
- Meta conversation fees scale with volume: utility template messages ~■15 each, marketing ~■70 each
- Up to 1,000 free service conversations/mo means most clients pay near-zero Meta fees

Per-client fixed via Meta direct: ~■6,000/mo (Resend + Mongo only).

Strategic implication: offer clients on cheaper plans (Starter) Meta direct, while higher-tier (Growth/Scale) get Unipile for hosted-auth convenience. Cuts Starter fixed cost by ~75%.

5 · Platform fixed costs (NOT per client)

These don't attribute per-client, but they need to be paid by *somebody's* margin.

Line item	■/mo	Notes
Railway hosting	~■16,000 (\$10)	Single service, scales modestly
Mongo Atlas M0 (free)	~■0 today	On free tier; switches to M10 (~■91k) around 10-12 clients
Resend Pro	~■32,000 (\$20)	Shared (counted partially above)
PostHog	0	Free tier sufficient until 1M events/mo

Line item	■/mo	Notes
OpenAI Whisper budget cap	~■240,000 (\$150)	Pre-funded monthly ceiling for voice transcription. Bursts above this require operator approval. Sized for ~25k minutes of inbound audio/mo (covers a full Tier-0 + Tier-1 cohort with comfortable headroom).
Anthropic Claude API budget cap	~■160,000 (\$100)	Pre-funded monthly ceiling for production drafts (Haiku 4.5 + occasional Vision OCR). Sized for ~3.5M tokens/mo at current per-token rates. Caller-side rate limits in <code>services/usage_meter.py</code> enforce the ceiling per client.
Claude Pro subscription (founder tooling)	~■32,000 (\$20)	The Pro subscription Yori uses to build, draft, plan, debug. Powers everything Claude Code touches. Upgrades to Max 5× (~■160k) likely once parallel sessions matter, Max 20× (~■320k) once team grows.
Domain + SSL	~■5,000/mo amortised	reachng.ng renewal
Platform fixed total (today)	~■485,000/mo	At 10 clients = ~■48,500/client. At 25 clients = ~■19,400/client. Jumps with Mongo M10 (~+■91k) around 10-12 clients.

Claude Pro → Max trajectory:

- 0-15 clients: Pro is fine (~■32k/mo)
- 15-50 clients: Max 5× (~■160k/mo) — operator runs parallel Claude Code sessions, builds faster
- 50+ clients: Max 20× (~■320k/mo) — multiple ops team members on Claude
- Even at Max 20× ÷ 100 clients = ■3,200/client overhead. Negligible at scale.

API budget caps — why pre-fund vs metered:

- Predictable line items investors can underwrite. No "AI costs went 4× this month" surprises.
- Hard ceilings on `usage_meter.py` protect runway during a viral spike or scrape attack.
- If we routinely hit the ceiling, that's a *good* problem — raise the cap and reprice.

Allocate platform fixed across active clients when computing true margin.

5.5 · Growth / marketing budget (separate from COGS)

This is **not** cost-of-goods-sold. It's the acquisition spend that compounds into paying clients. Tracked separately so unit economics stay honest.

Line item	■/mo	Notes
Meta Ads (Facebook + Instagram)	~■320,000 (\$200)	Lagos-targeted lookalike + retargeting campaigns. Vertical-specific creative per SCENE_PACK. CAC target ≤ 30% of first month's MRR.
Google Ads	~■240,000 (\$150)	Branded + 5-6 high-intent terms ("WhatsApp AI assistant Lagos", "auto-reply customers Nigeria"). Mostly defensive once SEO ranks.
SEO tooling (Ahrefs / SEMrush starter)	~■160,000 (\$99)	Keyword tracking, backlink audit, competitor monitoring (johahq + any new Lagos rivals).
LinkedIn Sales Navigator (founder outreach)	~■128,000 (\$79)	For Yori's direct outreach to luxury RE / legal / clinic founders.
Content production (occasional freelance)	~■200,000	One case-study write-up or video edit per month once first paid client lands.
Apify (SDR funnel discovery)	~■16,000 (\$10)	Our own internal lead-discovery for the Prospect OS funnel. Was previously double-counted under platform fixed.
Marketing total (today)	~■1,064,000/mo (~\$665)	Roughly \$0.66k/mo at Year 1 ramp. Scales to \$1.5-2k/mo by Month 6 if early signals are strong.

Acquisition math at this spend:

- Target blended CAC ≤ ■300,000 (~ 2× a Starter month) — every paid client recovers CAC inside 60 days.
- At 5 net-new clients/mo, blended CAC = ■213k. Within target.
- At 2 net-new clients/mo, blended CAC = ■532k. Outside target — pause Ads, lean into founder outreach + organic.
- Re-evaluate spend split monthly. Cut what doesn't pay back; double the channel that does.

6 · Total cost per client per month

Combining sections 3 + 4 + 5 (allocating platform fixed across 10 clients):

Profile	Per-call	Per-client fixed (Unipile)	Per-client fixed (Meta direct)	+ Platform share	TOTAL (Unipile)	TOTAL (Meta)
Light	2,520	30,160	6,160	15,200	~■47,880/mo	~■23,880/mo
Medium	7,620	30,160	6,160	15,200	~■52,980/mo	~■28,980/mo
Heavy	18,840	33,160	9,160	15,200	~■67,200/mo	~■43,200/mo

Margin at proposed pricing (current PRODUCTS.md ladder)

Plan	Price/mo	Light client (Unipile)	Light client (Meta)	Medium (Unipile)	Heavy (Unipile)
Starter ■80k	80,000	+■32,120 (40%)	+■56,120 (70%)	+■27,020 (34%)	+■12,800 (16%)
Growth ■150k	150,000	+■102,120 (68%)	+■126,120 (84%)	+■97,020 (65%)	+■82,800 (55%)
Scale ■300k	300,000	+■252,120 (84%)	+■276,120 (92%)	+■247,020 (82%)	+■232,800 (78%)

Read: **Starter via Unipile on a light client = 40% margin. Tight.** Same plan via Meta direct = 70% margin. **This is why we have to route Starter clients through Meta Cloud direct, not Unipile.**

Growth and Scale absorb Unipile fine. Heavy clients on Starter are unprofitable to dangerous (16% margin doesn't cover surprise spikes).

Decision matrix from this table

- **Move Starter to Meta Cloud API only.** No Unipile pairing on Starter. Healthy 70% margin.
- **Growth gets Unipile.** 65% margin on a medium client is the sweet spot.
- **Scale gets Unipile + priority + Co-Pilot extras.** 78–82% margin on heavy. Real money per client.
- **Cap Starter usage** to Light profile (~■2,500/mo per-call). If a client trends Medium on Starter, auto-prompt upgrade ("you're using Growth-level volume, here's what changes").

7 · The runaway / abuse cases

A single rogue client can torch margin. The meter has anti-runaway already; here's what protects us:

Feature	Per-minute hard cap (in code)	Effect of cap
voice	20 calls/min	Max ■9,600/hour on Whisper
receipt	20 calls/min	Max ■12,000/hour on vision

Feature	Per-minute hard cap (in code)	Effect of cap
classifier	60 calls/min	Max ■7,200/hour
drafter	60 calls/min	Max ■14,400/hour
memory	60 calls/min	Max ■10,800/hour
copilot	15 calls/min	Max ■7,200/hour

These caps are short-window; sustained abuse over hours is what the daily/monthly alert layer needs to catch. **Recommended addition:** flag any client whose 7-day rolling cost exceeds 2× the cost-of-tier baseline, surface on admin dashboard as `at_risk`.

8 · What the admin dashboard sees today

Already shipped (`templates/dashboard.html` line ~2400, `api/billing.py`):

- Per-client table: revenue, API cost MTD, margin %, top spend feature, `at_risk` flag
- Cohort totals: platform-wide revenue, cost, margin, client count
- Per-client drill-in via `/api/v1/admin/billing/{client_id}` showing usage by feature

What this view is missing (the gap to close before first paid client):

- **Per-client fixed costs not attributed.** The table shows API costs only. Unipile, Resend share, Mongo share are zero in the UI. Real margin is **always lower** than shown.
- **No platform-fixed allocation.** A "platform overhead" line per client isn't included.
- **No trend chart.** Just MTD totals. We need 30/60/90-day series to spot rising usage before it becomes a margin emergency.
- **No usage-by-profile classification.** We can compute whether a Starter client is exhibiting Medium/Heavy usage. Auto-prompt upgrade.
- **No runaway alerting.** `at_risk` flag exists but only logs; no Slack / WhatsApp / email push to the operator.

9 · Recommended dashboard upgrade (Phase next)

Order of priority for the build:

- **Attribute per-client fixed costs** in the billing query so margin % reflects reality. Add a `fixed_cost_ngn` field on `clients` doc (default ■24k Unipile / ■6k Meta) and add to the total. (1 hour)
- **Allocate platform-fixed** by dividing the manually-entered `platform_overhead_ngn` (current ■152k) across active client count. (30 min)
- **Profile classifier** that reads usage and labels each client `light` / `medium` / `heavy`. Surface on the row so we can spot upgrade prompts. (1 hour)
- **30-day trend chart** per client (cost line + revenue line). Tiny inline sparkline in the row. (2 hours)
- **Runaway push alert** when 7-day rolling cost exceeds 2× tier baseline → Slack webhook + owner WhatsApp via `OWNER_WHATSAPP` env. (1 hour)

Total: ~5.5 hours to make the dashboard production-grade for margin defense.

10 · The numbers to lock in the pricing session

When we sit down for pricing today, the only inputs you need to bring are:

- Your **target gross margin %** per tier (industry norm for SaaS: Starter 50%, Growth 70%, Scale 80%)
- Your **floor instinct** — what would feel wrong to charge less than? (Anchors the bottom)
- Your **ceiling instinct** — what would make a premium Lagos owner say "that's the price of a junior staffer's monthly salary, no thanks"? (Anchors the top)

Everything else (cost per profile, margin at price, breakeven volume) is in this doc.

My recommendation as a starting point:

- **Starter ■80k** → keep, route via Meta Cloud only. 70% margin on a light client. Healthy.
- **Growth ■150k** → keep. 65% margin on medium client via Unipile. Sweet spot.
- **Scale ■300k** → keep. 78%+ margin even on heavy clients.

Then tune from there based on your reads.